

AI Model Performance Categorization

The following table categorizes specific AI models based on benchmarked performance across complex reasoning, standard text-to-text generation, and multilingual linguistic tasks.

Reasoning	Text to Text	Multilingual
GPT OSS 120B	GPT OSS 120B	GPT OSS 120B
GPT OSS 20B	GPT OSS 20B	GPT OSS 20B
Qwen 3 32B	Llama 4 Scout	Llama 4 Scout
	Llama 3.3 70B	Llama 3.3 70B
		Whisper Large v3

Service Provider Metrics and Pricing Analysis

This analysis details current service provider offerings, standardizing operational costs per one million tokens and defining context window limitations for production environments.

Service Provider	Model	Context Window	Input Cost (per 1M tokens)	Output Cost (per 1M tokens)
Anthropic	Claude Opus 4.6	1M	\$5/M	\$25/M
Anthropic	Claude Sonnet 4.6	1M	\$3/M	\$15/M

Anthropic	Claude Haiku 4.5	200K	\$1/M	\$5/M
OpenAI	GPT-5.4 Nano	400K	\$0.20/M	\$1.25/M
OpenAI	GPT-5.4 Mini	400K	\$0.75/M	\$4.50/M
OpenAI	GPT-5.4 Pro	1.05M	\$30/M	\$180/M
Google Gemini	gemini-2.5-flash-lite	1.05M	\$0.10/M	\$0.40/M
Google Gemini	gemini-embedding-001	20K	\$0.15/M	\$0/M
Google Gemini	gemini-2.5-flash-image	33k	\$0.30/M	2.5/M

Tokenization Examples and Analysis

The following breakdown illustrates tokenization behavior using the OpenAI Tokenizer, providing granular data on token counts, character lengths, and specific numeric identifiers (Token IDs) for the GPT-4 series of models.

Tokenization Example 1

"How to create a login page using TailWind CSS?"

- **Token Count:** 11
- **Character Count:** 45
- **Token IDs:** [5299, 316, 2501, 261, 9464, 3011, 2360, 64498, 62417, 49194, 13]

Tokenization Example 2

"What does the word hapless mean?"

- **Token Count:** 8
- **Character Count:** 32
- **Token IDs:** [4827, 2226, 290, 2195, 35356, 2695, 4774, 30]

Tokenization Example 3

"Providing background information about the user's business, the specific project, or the scenario. For instance, a 'master prompt' might include an entire 30-page document of company-specific details for the AI to reference."

- **Token Count:** 42
- **Character Count:** 224
- **Token IDs:** [146804, 6357, 2164, 1078, 290, 49366, 2413, 11, 290, 4857, 2993, 11, 503, 290, 25697, 13, 2214, 4934, 11, 261, 392, 22522, 15226, 1, 3572, 3931, 448, 6508, 220, 1130, 29704, 3213, 328, 3175, 31999, 4878, 395, 290, 20837, 316, 9682, 13]